

ml_world_population

January 25, 2023

1 Step 1: Importing libraries

import the necessary libraries to work with the dataset and the machine learning models.

```
[ ]: # Import the necessary libraries

import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split

import warnings
warnings.filterwarnings('ignore') # to ignoring warnings
```

2 Step 2: Reading the dataset

```
[ ]: # We have imported dataset of world

df = pd.read_csv('world_population.csv')
```

3 Step 3: EDA/Wrangling

Check the data type and other for EDA and Wrangling

```
[ ]: # check the data for understanding of data
df
```

```
[ ]:
```

	Rank	CCA3	Country/Territory	Capital	Continent	\
0	36	AFG	Afghanistan	Kabul	Asia	
1	138	ALB	Albania	Tirana	Europe	
2	34	DZA	Algeria	Algiers	Africa	
3	213	ASM	American Samoa	Pago Pago	Oceania	
4	203	AND	Andorra	Andorra la Vella	Europe	
..	
229	226	WLF	Wallis and Futuna	Mata-Utu	Oceania	
230	172	ESH	Western Sahara	El Aaiún	Africa	

231	46	YEM	Yemen	Sanaa	Asia
232	63	ZMB	Zambia	Lusaka	Africa
233	74	ZWE	Zimbabwe	Harare	Africa

	2022 Population	2020 Population	2015 Population	2010 Population	\
0	41128771	38972230	33753499	28189672	
1	2842321	2866849	2882481	2913399	
2	44903225	43451666	39543154	35856344	
3	44273	46189	51368	54849	
4	79824	77700	71746	71519	
..	
229	11572	11655	12182	13142	
230	575986	556048	491824	413296	
231	33696614	32284046	28516545	24743946	
232	20017675	18927715	16248230	13792086	
233	16320537	15669666	14154937	12839771	

	2000 Population	1990 Population	1980 Population	1970 Population	\
0	19542982	10694796	12486631	10752971	
1	3182021	3295066	2941651	2324731	
2	30774621	25518074	18739378	13795915	
3	58230	47818	32886	27075	
4	66097	53569	35611	19860	
..	
229	14723	13454	11315	9377	
230	270375	178529	116775	76371	
231	18628700	13375121	9204938	6843607	
232	9891136	7686401	5720438	4281671	
233	11834676	10113893	7049926	5202918	

	Area (km ²)	Density (per km ²)	Growth Rate	World Population Percentage
0	652230	63.0587	1.0257	0.52
1	28748	98.8702	0.9957	0.04
2	2381741	18.8531	1.0164	0.56
3	199	222.4774	0.9831	0.00
4	468	170.5641	1.0100	0.00
..
229	142	81.4930	0.9953	0.00
230	266000	2.1654	1.0184	0.01
231	527968	63.8232	1.0217	0.42
232	752612	26.5976	1.0280	0.25
233	390757	41.7665	1.0204	0.20

[234 rows x 17 columns]

```
[ ]: # Verified the data type and Null values
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 234 entries, 0 to 233
Data columns (total 17 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Rank                                  234 non-null    int64
1   CCA3                                  234 non-null    object
2   Country/Territory                    234 non-null    object
3   Capital                              234 non-null    object
4   Continent                            234 non-null    object
5   2022 Population                      234 non-null    int64
6   2020 Population                      234 non-null    int64
7   2015 Population                      234 non-null    int64
8   2010 Population                      234 non-null    int64
9   2000 Population                      234 non-null    int64
10  1990 Population                      234 non-null    int64
11  1980 Population                      234 non-null    int64
12  1970 Population                      234 non-null    int64
13  Area (km²)                          234 non-null    int64
14  Density (per km²)                   234 non-null    float64
15  Growth Rate                         234 non-null    float64
16  World Population Percentage          234 non-null    float64
dtypes: float64(3), int64(10), object(4)
memory usage: 31.2+ KB
```

```
[ ]: # Printed columns name for X and y
df.columns
```

```
[ ]: Index(['Rank', 'CCA3', 'Country/Territory', 'Capital', 'Continent',
          '2022 Population', '2020 Population', '2015 Population',
          '2010 Population', '2000 Population', '1990 Population',
          '1980 Population', '1970 Population', 'Area (km²)', 'Density (per km²)',
          'Growth Rate', 'World Population Percentage'],
          dtype='object')
```

4 Step 4: Defining the independent and dependent variables

```
[ ]: # defining variables for X and y
X = df[[
    '2022 Population', '2020 Population', '2015 Population',
    '2010 Population', '2000 Population', '1990 Population']]
a = df['World Population Percentage'] # y is replaced with a in this step and
↳ in next
```

5 Step 5: Transforming the dependent variable

```
[ ]: # We changed the values by encoding as the y is countinous
from sklearn import preprocessing
from sklearn import utils
lab = preprocessing.LabelEncoder()
y = lab.fit_transform(a)
```

6 Step 6: Splitting the dataset

```
[ ]: #Splited the dataset in two parts for test and train, we used the random state
    ↳to get \
# the same results each time, if we select none everytime results will be
    ↳changed
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
    ↳random_state=42)
```

7 Step 7: Import ML Libraries

Import the required libraries for the models

```
[ ]: # Imported required Libraries for Machine learning
from sklearn.linear_model import LogisticRegression
from sklearn.svm import SVC
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import accuracy_score, f1_score, precision_score,
    ↳recall_score
from sklearn.model_selection import train_test_split
```

8 Step 8: Calculating Accuracy score of each model

Selecting, Iterating over the models and Sorting the models to check accuracy sore of each Model

```
[ ]: # we selected the Machine leaning models that we can use
models = [LogisticRegression(), SVC(), DecisionTreeClassifier(),
    ↳RandomForestClassifier(), KNeighborsClassifier()]
model_names = ['Logistic Regression', 'SVM', 'Decision Tree', 'Random Forest',
    ↳'KNN']
# we loopthrough each model and save the accuracy score in model_name
models_scores = []
for model, model_name in zip(models, model_names):
    model.fit(X_train, y_train)
    y_pred = model.predict(X_test)
```

```

    accuracy = accuracy_score(y_test, y_pred)
    models_scores.append([model_name, accuracy])
# used the lambda function to loopthroug each score type and sort from a to z
sorted_models = sorted(models_scores, key=lambda x: x[1], reverse=True)
for model in sorted_models:
    print("Accuracy Score: ", f'{model[0]} : {model[1]:.2f}')

```

c:\Users\muham\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\linear_model_logistic.py:458: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
```

Accuracy Score: Decision Tree : 0.62

Accuracy Score: Random Forest : 0.57

Accuracy Score: KNN : 0.40

Accuracy Score: SVM : 0.15

Accuracy Score: Logistic Regression : 0.00

9 Step 9: Calculating precision score of each model

Selecting, Iterating over the models and Sorting the models to check precision_score of each Model

```

[ ]: models = [LogisticRegression(), SVC(), DecisionTreeClassifier(),
    ↪RandomForestClassifier(), KNeighborsClassifier()]
model_names = ['Logistic Regression', 'SVM', 'Decision Tree', 'Random Forest',
    ↪'KNN']
models_scores = []
for model, model_name in zip(models, model_names):
    model.fit(X_train, y_train)
    y_pred = model.predict(X_test)
    Precision = precision_score(y_test, y_pred, average='micro') # Included
    ↪average='micro' for calcluating averaging score of each value
    models_scores.append([model_name, Precision])

sorted_models = sorted(models_scores, key=lambda x: x[1], reverse=True)
for model in sorted_models:
    print("Precision Score: ", f'{model[0]} : {model[1]:.2f}')

```

c:\Users\muham\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\linear_model_logistic.py:458: ConvergenceWarning: lbfgs failed to converge (status=1):

STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
```

Precision Score: Decision Tree : 0.60

Precision Score: Random Forest : 0.60

Precision Score: KNN : 0.40

Precision Score: SVM : 0.15

Precision Score: Logistic Regression : 0.00

10 Step 10: Calculating Recall score of each model

Selecting, Iterating over the models and Sorting the models to check Recall_score of each Model

```
[ ]: models = [LogisticRegression(), SVC(), DecisionTreeClassifier(),  
               ↳ RandomForestClassifier(), KNeighborsClassifier()]  
model_names = ['Logistic Regression', 'SVM', 'Decision Tree', 'Random Forest',  
               ↳ 'KNN']  
models_scores = []  
for model, model_name in zip(models, model_names):  
    model.fit(X_train, y_train)  
    y_pred = model.predict(X_test)  
    Recall = recall_score(y_test, y_pred, average='micro')  
    models_scores.append([model_name, Recall])  
  
sorted_models = sorted(models_scores, key=lambda x: x[1], reverse=True)  
for model in sorted_models:  
    print("Precision Score: ", f'{model[0]} : {model[1]:.2f}')
```

c:\Users\muham\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\linear_model_logistic.py:458: ConvergenceWarning: lbfgs failed to converge (status=1):

STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
```

Precision Score: Random Forest : 0.62

Precision Score: Decision Tree : 0.60

Precision Score: KNN : 0.40

Precision Score: SVM : 0.15
Precision Score: Logistic Regression : 0.00

11 Step 11: Calculating F1 score of each model

Selecting, Iterating over the models and Sorting the models to check f1_score of each Model

```
[ ]: models = [LogisticRegression(), SVC(), DecisionTreeClassifier(),  
               ↳ RandomForestClassifier(), KNeighborsClassifier()]  
model_names = ['Logistic Regression', 'SVM', 'Decision Tree', 'Random Forest',  
               ↳ 'KNN']  
models_scores = []  
for model, model_name in zip(models, model_names):  
    model.fit(X_train, y_train)  
    y_pred = model.predict(X_test)  
    F1 = f1_score(y_test, y_pred, average='micro')  
    models_scores.append([model_name, F1])  
  
sorted_models = sorted(models_scores, key=lambda x: x[1], reverse=True)  
for model in sorted_models:  
    print("F1 Score: ", f'{model[0]} : {model[1]:.2f}')
```

c:\Users\muham\AppData\Local\Programs\Python\Python311\Lib\site-
packages\sklearn\linear_model_logistic.py:458: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear_model.html#logistic-
regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result(
```

F1 Score: Random Forest : 0.60
F1 Score: Decision Tree : 0.57
F1 Score: KNN : 0.40
F1 Score: SVM : 0.15
F1 Score: Logistic Regression : 0.00

12 Step 12: Perform grid search to check best score and best parameter

for best parameters for each model

```
[ ]: # Imported GridsearchCV
```

```

from sklearn.model_selection import GridSearchCV

models = [LogisticRegression(), SVC(), DecisionTreeClassifier(),
↳ RandomForestClassifier(), KNeighborsClassifier()]
model_names = ['Logistic Regression', 'SVM', 'Decision Tree', 'Random Forest',
↳ 'KNN']
models_scores = []

for model, model_name in zip(models, model_names):
    # Define the grid search parameters
    if model_name == 'Logistic Regression':
        params = {'C': [0.1, 1, 10]}
    elif model_name == 'SVM':
        params = {'C': [0.1, 1, 10], 'kernel': ['linear', 'rbf']}
    elif model_name == 'Decision Tree':
        params = {'max_depth': [1, 2, 3, 4, 5, 6, 7, 8, 9, 10]}
    elif model_name == 'Random Forest':
        params = {'n_estimators': [10, 50, 100], 'max_depth': [1, 2, 3, 4, 5,
↳ 6, 7, 8, 9, 10]}
    elif model_name == 'KNN':
        params = {'n_neighbors': [3, 5, 7, 9, 11, 13, 15]}
    else:
        params = {}
    # Perform the grid search
    grid_search = GridSearchCV(model, param_grid=params, cv=5)
    grid_search.fit(X_train, y_train)
    model = grid_search.best_estimator_
    y_pred = model.predict(X_test)
    Recall = recall_score(y_test, y_pred, average='micro')
    models_scores.append([model_name, Recall])

# Print best parameters and best score
print("Best Parameters: ", grid_search.best_params_)
print("Best Score: ", grid_search.best_score_)

```

Best Parameters: {'n_neighbors': 3}
Best Score: 0.5453769559032717

13 Step 13: Perform grid search to check best Model

```

[ ]: models = [LogisticRegression(), SVC(), DecisionTreeClassifier(),
↳ RandomForestClassifier(), KNeighborsClassifier()]
model_names = ['Logistic Regression', 'SVM', 'Decision Tree', 'Random Forest',
↳ 'KNN']

```



```

models_scores = []
best_model = None
best_score = 0

for model, model_name in zip(models, model_names):
    # Define the grid search parameters
    if model_name == 'Logistic Regression':
        params = {'C': [0.1, 1, 10]}
    elif model_name == 'SVM':
        params = {'C': [0.1, 1, 10], 'kernel': ['linear', 'rbf']}
    elif model_name == 'Decision Tree':
        params = {'max_depth': [1, 2, 3, 4, 5, 6, 7, 8, 9, 10]}
    elif model_name == 'Random Forest':
        params = {'n_estimators': [10, 50, 100], 'max_depth': [1, 2, 3, 4, 5, 6, 7, 8, 9, 10]}
    elif model_name == 'KNN':
        params = {'n_neighbors': [3, 5, 7, 9, 11, 13, 15]}
    else:
        params = {}
    # Perform the grid search
    grid_search = GridSearchCV(model, param_grid=params, cv=5)
    grid_search.fit(X_train, y_train)
    model = grid_search.best_estimator_
    y_pred = model.predict(X_test)
    Recall = recall_score(y_test, y_pred, average='micro')
    models_scores.append([model_name, Recall])
    if grid_search.best_score_ > best_score:
        best_score = grid_search.best_score_
        best_model = model

print("Best Parameters: ", grid_search.best_params_)
print("Best Score: ", grid_search.best_score_)
print(f'Best Model: {best_model}')

```

```

Best Parameters: {'n_neighbors': 3}
Best Score: 0.5453769559032717
Best Model: SVC(C=0.1, kernel='linear')

```